

Systematic Error Handling in C++

Prepared for The C++ and Beyond Seminar 2012

Andrei Alexandrescu, Ph.D.

`andrei@erdani.com`

Agenda

- Introduction
- Expected<T>
- ScopeGuard11

Introduction

What is Error Handling?

Means and techniques dedicated to handling possible but unexpected and useless inputs

Part of *Reliability*

- Error handling
 - working core
 - correct program
 - incorrect inputs
- General reliability
 - non-working core
 - incorrect program
 - incorrect inputs

Error Handling Scenarios

- User input errors
- File I/O errors
- Networking errors
- Environmental issues
- Device malfunctions
- ...

Reliability Scenarios

- Program bugs
- Unexpected program state
- Corrupted I/O
- Core memory failure
- . . .

Hypothesis

- Bad error handling engenders errors

Hypothesis

- Bad error handling engenders errors
- Insufficient testing makes improbable error scenarios never covered

Hypothesis

- Bad error handling engenders errors
- Insufficient testing makes improbable error scenarios never covered
- Legit errors lead to corrupt program

Hypothesis

- Bad error handling engenders errors
- Insufficient testing makes improbable error scenarios never covered
- Legit errors lead to corrupt program

Example: Thunderbird

- Handling of networking errors almost rigorously bad
- Modal error dialogs
- Numerous race conditions
- Metastable states
- Broken core invariants

Expected **Types**

Motivation

- Exceptions are, um, good
- Slow on the exceptional path
- Hopelessly serial
 - Only one exception in flight
 - Requires immediate, exclusive attention
 - Dedicated control flow
- Associated only with root reasons, not goals
 - “I/O error” doesn’t describe “saving weight file”

Key Insight

- Common complaint: “Error codes are limited! Exceptions are arbitrarily rich!”
- Make exceptions *be* the error codes!

Key Idea

Expected<T> is either a T or the exception preventing its creation

Instead of this...

// Returns 0 on error and sets errno. Ahem.
`int` parseInt(`const` string&);

... or this...

```
// Throws invalid_input or overflow  
int parseInt(const string&);
```

... have this!

```
// Returns an expected int  
Expected<int> parseInt(const string&);
```

Related Work

- Haskell: `Maybe T`
 - Scala: `Option[T]`
 - C#: `Nullable<T>`
 - Boost: `optional<T>`
-
- Either a value of type `T`, or no value at all
 - Defines primitives to store, test, extract
 - Not packed with an exception

Related Work

- C++11: `promise<T>/future<T>`
- Either a value of type `T`, or an exception
- Primitives focused on inter-thread, async communication
- We want eager, synchronous

Expected<T> characteristics

- Associates errors with computational goals
- Naturally allows multiple exceptions in flight
- Switch between “error handling” and “exception throwing” styles
- Teleportation possible
 - Across thread boundaries
 - Across **nothrow** subsystem boundaries
 - Across time: save now, throw later
- Collect, group, combine exceptions

Expected<T> implementation

- Basic idea: Variant with a T or a `std::exception_ptr`

```
template <class T> class Expected {  
    union {  
        T ham;  
        std::exception_ptr spam;  
    };  
    bool gotHam;  
    Expected() {} // used internally  
public:
```

Expected<T> implementation

```
Expected(const T& rhs) : ham(rhs), gotHam(true) {}
Expected(T&& rhs)
    : ham(std::move(rhs))
    , gotHam(true) {}
Expected(const Expected& rhs) : gotHam(rhs.gotHam) {
    if (gotHam) new(&ham) T(rhs.ham);
    else new(&spam) std::exception_ptr(rhs.spam);
}
Expected(Expected&& rhs) : gotHam(rhs.gotHam) {
    if (gotHam) new(&ham) T(std::move(rhs.ham));
    else new(&spam)
        std::exception_ptr(std::move(rhs.spam));
}
```



```
void swap(Expected& rhs) {  
    if (gotHam) {  
        if (rhs.gotHam) {  
            using std::swap;  
            swap(ham, rhs.ham);  
        } else {  
            auto t = std::move(rhs.spam);  
            new(&rhs.ham) T(std::move(ham));  
            new(&spam) std::exception_ptr(t);  
            std::swap(gotHam, rhs.gotHam);  
        }  
    } else {  
        if (rhs.gotHam) {  
            rhs.swap(*this);  
        } else {  
            spam.swap(rhs.spam);  
            std::swap(gotHam, rhs.gotHam);  
        }  
    }  
}
```

Building from exception

```
template <class E>
static Expected<T> fromException(const E& exception) {
    if (typeid(exception) != typeid(E)) {
        throw std::invalid_argument(
            "slicing detected");
    }
    return fromException(
        std::make_exception_ptr(exception));
}
```

Building from exception

```
static Expected<T> fromException(std::exception_ptr p) {  
    Expected<T> result;  
    result.gotHam = false;  
    new(&result.spam) std::exception_ptr(std::move(p));  
    return result;  
}
```

```
static Expected<T> fromException() {  
    return fromException(std::current_exception());  
}
```

Access

```
bool valid() const {  
    return gotHam;  
}
```

```
T& get() {  
    if (!gotHam) std::rethrow_exception(spam);  
    return ham;  
}
```

```
const T& get() const {  
    if (!gotHam) std::rethrow_exception(spam);  
    return ham;  
}
```

Probing for Type

```
template <class E>
bool hasException() const {
    try {
        if (!gotHam) std::rethrow_exception(spam);
    } catch (const E& object) {
        return true;
    } catch (...) {
    }
    return false;
}
```

Icing

```
template <class F>
static Expected fromCode(F fun) {
    try {
        return Expected(fun());
    } catch (...) {
        return fromException();
    }
}

auto r = Expected<string>::fromCode([] {
    ...
});
```

Use example (callee)

```
Expected<int> parseInt(const std::string& s) {  
    int result;  
    ...  
    if (nonDigit) {  
        return Expected<int>::fromException(  
            std::invalid_argument("not a number"));  
    }  
    ...  
    if (tooManyDigits) {  
        return Expected<int>::fromException(  
            std::out_of_range("overflow"));  
    }  
    ...  
    return result;  
}
```

Use example (caller)

```
// Caller
string s = readline();
auto x = parseInt(s).get(); // throw on error
auto y = parseInt(s); // won't throw
if (!y.valid()) {
    // handle locally
    if (y.hasException<std::invalid_argument>()) {
        // no digits
        ...
    }
    y.get(); // just "re"throw
}
```


Expected<T> aftermath

- Encodes a value OR the error in attempting to produce it
- Groups data with error flow
- Non-serial in nature
- Supports different styles of coding
- Arbitrarily rich “error codes”

ScopeGuard11

Recall ScopeGuard

- Introduced in 2000 article
- Makes transactional code linear
- Lightweight RAI
- Scales well to multi-step transactions

⟨action⟩

⟨action⟩

⟨cleanup⟩

⟨action⟩

⟨cleanup⟩
⟨next⟩

⟨action⟩

⟨cleanup⟩

⟨next⟩

⟨rollback⟩

C

```
if (<action>) {  
    if (!<next>) {  
        <rollback>  
    }  
    <cleanup>  
}
```


C+

```
class RAII {  
    RAII() { <action> }  
    ~RAII() { <cleanup> }  
};  
  
...  
RAII raii;  
try {  
    <next>  
} catch (...) {  
    <rollback>  
    throw;  
}
```

Entering Composition

C (expand $\langle \text{next}_1 \rangle$)

```
if ( $\langle \text{action}_1 \rangle$ ) {  
    if ( $\langle \text{action}_2 \rangle$ ) {  
        if (! $\langle \text{next}_2 \rangle$ ) {  
             $\langle \text{rollback}_2 \rangle$   
             $\langle \text{rollback}_1 \rangle$   
        }  
         $\langle \text{cleanup}_2 \rangle$   
    } else {  
         $\langle \text{rollback}_1 \rangle$   
    }  
     $\langle \text{cleanup}_1 \rangle$   
}
```

C+

```
class RAII1 {  
    RAII1() { <action1> }  
    ~RAII1() { <cleanup1> }  
};  
class RAII2 {  
    RAII2() { <action2> }  
    ~RAII2() { <cleanup2> }  
};  
...
```

C+ (expand $\langle \text{next}_1 \rangle$)

```
RAII1 raii1;
try {
    RAII2 raii2;
    try {
         $\langle \text{next}_2 \rangle$ 
    } catch (...) {
         $\langle \text{rollback}_2 \rangle$ 
        throw;
    }
} catch (...) {
     $\langle \text{rollback}_1 \rangle$ 
    throw;
}
```

Dislocation + Nesting =
Fail

C++11 with ScopeGuard

⟨action⟩

```
auto g1 = scopeGuard([] { ⟨cleanup⟩ });
```

```
auto g2 = scopeGuard([] { ⟨rollback⟩ });
```

⟨next⟩

```
g2.dismiss();
```

ScopeGuard composition

$\langle \text{action}_1 \rangle$

```
auto g1 = scopeGuard([&] {  $\langle \text{cleanup}_1 \rangle$  });
```

```
auto g2 = scopeGuard([&] {  $\langle \text{rollback}_1 \rangle$  });
```

$\langle \text{action}_2 \rangle$

```
auto g3 = scopeGuard([&] {  $\langle \text{cleanup}_2 \rangle$  });
```

```
auto g4 = scopeGuard([&] {  $\langle \text{rollback}_2 \rangle$  });
```

$\langle \text{next}_2 \rangle$

```
g2.dismiss();
```

```
g4.dismiss();
```


Macro Edition

```
⟨action1⟩  
SCOPE_EXIT { ⟨cleanup1⟩ };  
auto g1 = scopeGuard([&] { ⟨rollback1⟩ });  
⟨action2⟩  
SCOPE_EXIT { ⟨cleanup2⟩ };  
auto g2 = scopeGuard([&] { ⟨rollback2⟩ });  
⟨next2⟩  
g2.dismiss();  
g1.dismiss();
```

Painfully Close to Ideal!

```
⟨action1⟩  
SCOPE_EXIT { ⟨cleanup1⟩ };  
SCOPE_FAIL { ⟨rollback1⟩ }; // nope  
⟨action2⟩  
SCOPE_EXIT { ⟨cleanup2⟩ };  
SCOPE_FAIL { ⟨rollback2⟩ }; // nope  
⟨next2⟩
```

How does ScopeGuard work?

- C++98: elaborate implementation featuring type erasure
- In C++11 we can exploit
 - Type inference and **auto**
 - No need for type erasure
 - Lambda functions
 - Defer arbitrary code
 - Move semantics
 - No spurious double cleanup

Implementation

```
template <class Fun>
class ScopeGuard {
    Fun f_;
    bool active_;
public:
    ScopeGuard(Fun f)
        : f_(std::move(f))
        , active_(true) {
    }
    ~ScopeGuard() { if (active_) f_(); }
    void dismiss() { active_ = false; }
```

... to be continued ...

Implementation (cont'd)

...continued ...

```
ScopeGuard() = delete;
ScopeGuard(const ScopeGuard&) = delete;
ScopeGuard& operator=(const ScopeGuard&) = delete;
ScopeGuard(ScopeGuard&& rhs)
    : f_(std::move(rhs.f_))
    , active_(rhs.active_) {
    rhs.dismiss();
}
};
```

Type Deduction

```
template <class Fun>
ScopeGuard<Fun> scopeGuard(Fun f) {
    return ScopeGuard<Fun>(std::move(f));
}
```

Use

```
void fun() {  
    char name[] = "/tmp/deleteme.XXXXXX";  
    auto fd = mkstemp(name);  
    auto g1 = scopeGuard([] {  
        fclose(fd);  
        unlink(name);  
    });  
    auto buf = malloc(1024 * 1024);  
    auto g2 = scopeGuard([] { free(buf); });  
  
    ... use fd and buf ...  
}
```

Pseudo-Statement

```
namespace detail {  
    enum class ScopeGuardOnExit {};  
    template <typename Fun>  
    ScopeGuard<Fun>  
    operator+(ScopeGuardOnExit, Fun&& fn) {  
        return ScopeGuard<Fun>(std::forward<Fun>(fn));  
    }  
}
```

```
#define SCOPE_EXIT \  
    auto ANONYMOUS_VARIABLE(SCOPE_EXIT_STATE) \  
    = ::detail::ScopeGuardOnExit() + [&]()
```


Preprocessor Wonders

```
#define CONCATENATE_IMPL(s1, s2) s1##s2
#define CONCATENATE(s1, s2) CONCATENATE_IMPL(s1, s2)

#ifdef __COUNTER__
#define ANONYMOUS_VARIABLE(str) \
    CONCATENATE(str, __COUNTER__)
#else
#define ANONYMOUS_VARIABLE(str) \
    CONCATENATE(str, __LINE__)
#endif
```

Lo and Behold

```
void fun() {  
    char name[] = "/tmp/deleteme.XXXXXX";  
    auto fd = mkstemp(name);  
    SCOPE_EXIT { fclose(fd); unlink(name); };  
    auto buf = malloc(1024 * 1024);  
    SCOPE_EXIT { free(buf); };  
  
    ... use fd and buf ...  
}
```

(if no “;” after lambda, error message is meh)

Summary

Summary

- Make error handling simple and systematic
- `ExpectedT` encapsulates values with their history of failure
- `ScopeGuard11` encapsulates scoped control flow